

A Framework for Heuristic Scheduling for Parallel Processing on Multicore Architecture: A Case Study with Multiview Video Coding

Yi Pang, *Student Member, IEEE*, Lifeng Sun, *Member, IEEE*, Jiangtao (Gene) Wen, *Senior Member, IEEE*, Fengyan Zhang, Weidong Hu, Wei Feng, and Shiqiang Yang, *Senior Member, IEEE*

Abstract—In this paper, using the Intel multicore architectures and the emerging multiview video coding standard, we introduce a framework for performing analysis, simulation, and evaluation of heuristics scheduling algorithms for implementing computationally intensive algorithms on multicore processors. The framework allows for accurate and quantitative characterization of the performance of dynamic scheduling algorithms for multimedia applications on different multicore processors without actual implementation of the scheduling algorithm and application on the actual platform. Experimental results demonstrate the effectiveness and scalability of our framework.

Index Terms—Directed acyclic graph, multicore architecture, scheduling algorithms.

I. INTRODUCTION

IN SPITE OF the rapid improvement of microprocessor technology, many applications remain outside of the realm of practicality on even the most cutting-edge microprocessor platforms. One such computationally intensive application is video encoding, which has evolved in the last two decades from MPEG-2 to advanced video coding (AVC), as well as multiview video coding (MVC) and scalable video coding. Compared with MPEG-2, currently the most widely deployed video coding standard in the world, AVC delivers twice the video coding efficiency at significantly higher computation and storage complexity, and much higher encoding processing bandwidth requirements. On the other hand, the ever-widening application of media processing and video coding technologies necessitates encoding and processing of increasingly higher resolution inputs. As a result, relative to the demands on processor and I/O performance, processor capabilities, albeit

Manuscript received January 30, 2009; revised May 22, 2009. First version published September 1, 2009; current version published October 30, 2009. This work was supported by the 973 Program under Grant 2006CB303103, the 863 Program under Grant 2006AA01Z321, and the Key 863 Program under Grant 2009AA01Z328. This paper was recommended by Associate Editor M. Mattavelli.

The authors are with the Department of Computer Science and Technology, Tsinghua University, Beijing 100084, China (e-mail: pangy@mails.tsinghua.edu.cn; sunlf@tsinghua.edu.cn; jtwen@tsinghua.edu.cn; ppsnow@gmail.com; mail@hwd.name; allenfw@gmail.com; yangshq@tsinghua.edu.cn).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/TCSVT.2009.2031463

advancing rapidly, remains lagging behind the ever-higher requirements. This situation is illustrated in Fig. 1, which shows the rate of increase for the relative computation complexity of different video coding standards, as well as the rate of growth in computational capabilities for unicore and multicore microprocessors over the same period of time. The rates of growth in both cases are normalized to the year 1995, when MPEG-2 was first widely adopted.

A recent trend in microprocessor design is to put multiple processing cores on the same processor. The presence of multiple cores not only improves the amount of data that the processor can handle in parallel, but also reduces the cost of data communications and the bandwidth required for data exchange prior to/after processing between the cores, which in turn reduces bandwidth requirement for the communication between the processor and off-chip memory, as well the power consumption of the entire chip. Asanovic *et al.* [1] expresses the landscape of parallel computing research.

The challenge, however, in utilizing multicore technology for computation and data intensive applications such as video coding, is to “map” the application onto a multicore processor, so that the cores are fully utilized and unnecessary data exchanges between the cores and especially between the processor and external, off-chip memory are eliminated to the greatest extent.

Most computation and data intensive applications such as media processing, involves computation intensive tasks that can be parallelized. The frequency and scope of these operations, however, are usually context dependent, which makes it difficult for a static scheduling algorithm to achieve optimized performance partitioning such applications onto multicore processors. As an example, in the baseline and main profiles of the AVC standard, the motion prediction residual of each 16×16 macroblock (MB) is partitioned into nonoverlapping 4×4 blocks that subsequently undergo 4×4 integer transform, quantization, and entropy coding. Although the transform of each 4×4 block is independent of other 4×4 blocks, the entropy coding of the coefficients after quantization is context dependent, and must be processed sequentially. Therefore, even if the transforms of the 4×4 blocks in an MB is partitioned onto different processing cores, the resultant coefficients must still be exchanged between the cores so that entropy coding could be carried out correctly. In

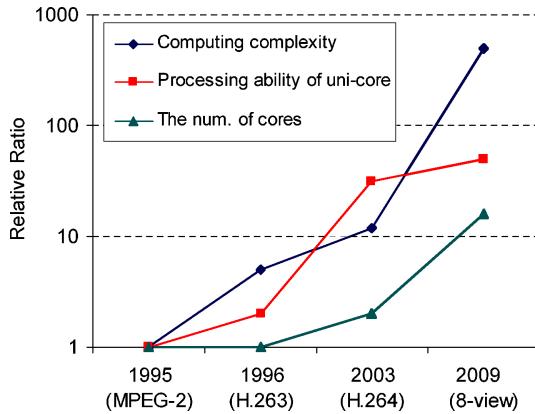


Fig. 1. Development of video compression and processor technologies.

addition, applications such as video coding present numerous possibilities of parallelization at different levels of the syntax hierarchy, resulting in different performance/complexity trade-offs. For example, in AVC encoding, a frame can be divided into multiple slices that can be independently encoded. Such mutually independent slices can, therefore, be processed on separate cores with minimal data exchanges between the cores during encoding. This approach may work nicely for processors with a small number (e.g., 2) of cores, especially when the input resolution is high. However, for processors with a large number (e.g., 8 or 16) of cores, dividing input frames into such a large number of slices leads to significant losses in coding efficiency that renders such partitioning useless.

In summary, the challenges in designing an application specific, good scheduling algorithm for multicore architectures lie in the following aspects: 1) the dependencies between the tasks in a specific application is usually context, implementation, and some times input dependent, making it difficult to design a “one-size-fits-all” scheduling algorithm; 2) evaluating different scheduling algorithms requires detailed real-time tracing information that are usually not available on the application layer; and 3) evaluating optimized scheduling algorithms requires trading off processor performance and application performance, for which a scientific framework of quantitative comparison has not been established. As a result, traditionally, virtually the only reliable way of evaluating the performance of a dynamic scheduling algorithm for multimedia applications on multicore processors is to implement the algorithm on the actual processor, along with the multimedia application itself, and the measure the performance using a large set of inputs. This approach is of course prohibitively expensive and is almost impossible to automate. Therefore, it is highly desirable to come up with an automatic tool that is capable of capturing the key characteristics of different designs (speed, bandwidth, utility, etc.) with minimal human intervention, so that automatic “tuning” of the design can be performed.

In this paper, we use the emerging MVC standard as an example and present a generic directed acyclic graph (DAG)-based framework that can be used to model the dynamic nature

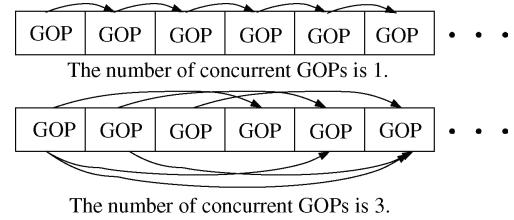


Fig. 2. Dependencies among GOPs as a function of NoC GOPs.

of the dependencies between tasks in a complicated application at different granularities. We demonstrate ways of using such a framework to come up with different heuristics for designing efficient scheduling algorithms, and more importantly, ways of evaluating or predicting the performance of scheduling algorithms for applications on multicore platforms without actually implementing the algorithm and the application on the target platform. In the paper, we use five different and commonly adopted heuristic-based scheduling algorithms for MVC with different encoding options as test cases. Through both analysis and experimental results, we show that the framework, tools, and methodology presented in this paper is capable of capturing the key characteristics and performance parameters of a scheduling algorithm.

The paper is organized as follows. Section II contains an analysis of the problem of scheduling video encoding tasks on multicore architectures. In Section III, we introduce the DAG-based framework, as well as methodologies for using such a framework for the design of heuristic for finding good scheduling algorithms and to analyze their performances. Using MVC as an illustrative example, in Section IV, we demonstrate the framework in operation and its usefulness, as well as experimental results. Finally, Section V contains the conclusions and suggestions for future research.

II. PARALLEL SCHEDULING OF VIDEO ENCODING FOR MULTICORE ARCHITECTURE

Video encoding is well known to be computationally demanding, and therefore, parallel architecture for video encoding has been an active research topic for well over a decade, and will continue to be so in the foreseeable future. Interested readers may refer to references [2]–[5]. In [6], Wang *et al.* proposed a novel high-definition television video decoder and decentralized control scheme with a data-driven architecture on the function level. The data-driven architecture was adapted to allow each processing unit to operate as soon as data and buffer become available. In [7], Nanda *et al.* introduced IBM’s video surveillance server prototype, implementing AVC on the cell broadband engine (Cell/B.E.) processor. The tasks are partitioned into four modules, each of which is assigned to a dedicated synergistic processing element.

In addition to single view coding, MVC as an extension to AVC was recently standardized to efficiently compress multiple parallel video inputs that capture different views of the same content. Although it is intuitive to see that MVC exhibits good inherent parallelism, tackling the encoding

and decoding of multiple views, requires higher processing power, more storage, and higher bandwidth communications. As a result, implementing MVC on multicore processors has become an important emerging research topic. In [8], a parallelization methodology for MVC based on hyper-space theory was presented and tested on multiprocessor platforms. [9] analyzed the parallelism of MVC using video frames as the parallel granularity. The techniques and methodology in these papers differ significantly as a result of the differences in target applications and platforms, which highlighted the need for a systematic and generic approach for performing theoretical analysis and design of scheduling and parallelization algorithms for video encoding on multicore processors.

A fundamental tool that we are using throughout this paper is DAG. Scheduling algorithm design and analysis for parallel systems and architectures using DAGs is a classic problem (see for example, [10]–[12]). This prior research, however, was not targeted at scheduling dynamic tasks on multicore architectures, but toward cluster, grid, and multiprocessor systems that involves tradeoffs that are different than MVC. These existing approaches targeted cases in which the DAGs are to a very large extent static, as opposed to involving dynamic insertion and deletion of nodes over time. Further more, the application of the DAG-based approach is mainly to facilitate only the designing of scheduling algorithms, as opposed to designing and then evaluating such algorithms, without actually implementing the scheduling algorithm in question.

In this paper, without losing generality, we use the total execution time as the criterion for measuring system performance, although other criteria are also possible.

In traditional sequential implementations of video encoding on a uncore processor, at each given moment, only one group of pictures (GOPs) can be processed and the encoding of the next GOP can not start until the current GOP is completed. In MVC and on multicore platforms, however, GOPs can be processed in parallel by the different cores. The number of GOPs that a multicore processor can process simultaneously is called the number of concurrent GOPs (NoC GOPs), which is limited by the delay constraints of the application and the sizes of the encoding and decoding buffers. When a processing core of the system is done with encoding a GOP of the input, it will “bypass” the other GOPs that the other cores are processing, and may go directly to the next GOP that has not been processed by any of the cores. In practice and to simplify data access strides, the allocation of the GOPs to be encoded by each core is usually periodic, as shown in Fig. 2 for the cases when the NoC GOPs is 1 or 3. The arrows in the figure illustrate the time order in which the GOPs are processed, and the interleaving of arrows depict the concurrent nature of the allocation of GOPs to cores.

Because the data dependencies in video encoding in general and MVC in particular, are complex, time variant, and sometimes dependent on the “earlier” decisions made by a dynamic scheduling algorithm, finding the optimal scheduling algorithm for MVC with a certain configurations (e.g., number of views, resolution, bit-rate, content, etc.) is often nondeterministic polynomial complete or at least immensely computational intensive, thereby making it difficult for embedding into real-

time encoding operations for improving system performance “on-the-fly.”

III. HEURISTIC PARALLEL SCHEDULING FRAMEWORK OF VIDEO ENCODING ON MULTICORE ARCHITECTURE

A. Problem Formulation

To model video encoding on a certain multicore processor, we first introduce a set of variables A which uniquely identifies a configuration of a particular encoding application on a particular multicore processor, e.g., MVC with eight views and other encoding parameters on IBM Cell/B.E. with eight cores, etc. The subset of A that is related to video encoding and impacts compression efficiency (Y) is denoted by A_{ve} . Obviously, Y is a function of A_{ve} , i.e., $Y = Y(A_{ve})$.

A particular video encoding application is characterized by a series of encoding tasks that may change over time. Depending on the particular application, the tasks could be defined at different granularities, e.g., encoding a GOP, a frame, a slice, so on, and so forth. When the i th task is processed, the next scheduling action $f(i)$ is constrained by the application and the platform, as well as earlier scheduling decisions. We denote the time between the completion of the $i - 1$ th task and the i th task by $\Delta t_j(i)$. It is straightforward to see that the total execution time $t_j(i)$ from beginning encoding to the completion of the i th task is simply $t_j(i) = \sum_{k=0}^i \Delta t_j(k)$. Given set A and a scheduling algorithm, the total execution time $t_j(n - 1)$ is uniquely decided, as is $T(n - 1) = \{t_j(n - 1), j = 0, 1, \dots, |Q| - 1\}$ where, n is the number of tasks, T is the set of all possible total execution times and Q is the set of heuristics scheduling algorithms. Usually, different heuristics lead to different scheduling algorithms. The problem of finding the optimal scheduling algorithm can be formulated as

$$\begin{aligned} & \min_{\{Q,A\}} T(n - 1) \\ &= \min_{\{Q,A\}} \{t_j = \sum_{i=0}^{n-1} \Delta t_j(i), j = 0, 1, \dots, |Q| - 1\} \\ & \text{s.t. } Y(A_{ve}) \geq Y_0 \end{aligned} \quad (1)$$

where Y_0 is the compression performance requirement of the application (e.g., as defined by target bit-rate/peak signal-to-noise ratio).

In the above formulation, the heuristics scheduling algorithm that leads to the shortest total execution time is considered optimal. As noted earlier, it is straightforward to generalize the formulation to other optimality criteria.

It should be noted that we assume that the video encoder is causal and 1-pass, i.e., it is not possible to revert to a decision that the encoder and/or the scheduler has made earlier. This is a realistic assumption for most real-time video applications

B. Application to MVC

There are different levels of the syntax elements in MVC that can be used as parallel granularities, including: GOP, frame, slice, MB, etc. MVC also consists of different encoding steps, including intra-prediction (Intra-P), inter-prediction

TABLE I
CONSTRAINTS ON THE GOP, FRAME, AND MB LEVELS

Parallel Granularity	Constraints Generated From		
	GOP Level	Frame Level	MB Level
MB	Concurrent GOP	Inter-P	Intra-P, ME, DCT, etc.
Frame	Concurrent GOP	Inter-P	Null
GOP	Concurrent GOP	Null	Null

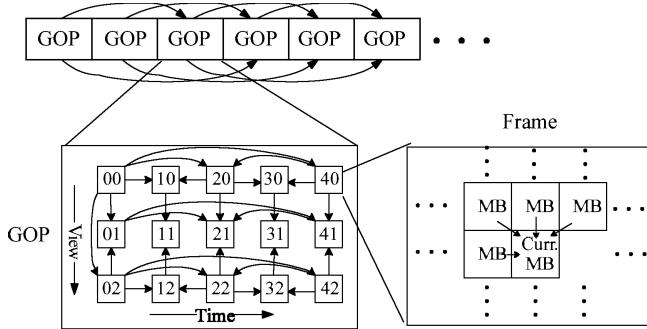


Fig. 3. Hierarchical dependence of MVC—an example.

(Inter-P), motion estimation (ME), discrete cosine transform (DCT), etc. The different constraints and tradeoff involved in selecting different syntax units as parallel granularities are summarized in Table I. For example, when we choose frame as parallel granularity (as the “frame” row in Table I), the dependencies between frames generate from GOP level and frame level constraints. MBs are within frame, so the constraints on MB level within a frame, e.g., Intra-P, ME, etc., do not impact dependencies on frame level (“null” in Table I). On GOP level, the constraint is concurrent GOP; on the frame level, the constraint is Inter-P which impacts dependencies among frames but within a GOP.

As an example, consider MVC with 12 frames in each GOP, 3 views and an NoC GOP of 2. The corresponding hierarchical prediction/dependency structure is depicted in Fig. 3.

The dependency of Fig. 3 can be modeled by a DAG. In the DAG, the notes designate the tasks $\tau_0, \tau_1, \dots, \tau_{n-1} \in \Gamma$, where Γ is the set of tasks. These nodes are weighted by their corresponding workloads $w_0, w_1, \dots, w_{n-1} \in W$. The dependencies between the tasks correspond to the edges in the DAG, namely $e_i = (\tau_j, \tau_k) \in E, \tau_j, \tau_k \in \Gamma$. The edges are in turn weighted by the communication costs of $h_0, h_1, \dots, h_{|E|-1} \in H$. By definition, τ_j depends on τ_k , if τ_j cannot start until τ_k finishes. Fig. 4 shows the DAG for MVC on GOP level, with NoC GOPs of 2.

In addition to the variables in A that are related to video encoding as A_{ve} , the remaining components that are not directly related to video encoding, but related to the platform and processor are denoted as A_{mc} . A_{mc} is mutually exclusive with A_{ve} . Processing capability of the cores and memory are denoted as nodes, and the directions and bandwidth of data transmission between cores and memory are weighted edges. They are all the elements in A_{mc} . Then a given multicore architecture can be also modeled as a weighted directed graph (DG). This DG is different from the DAG modeling the MVC

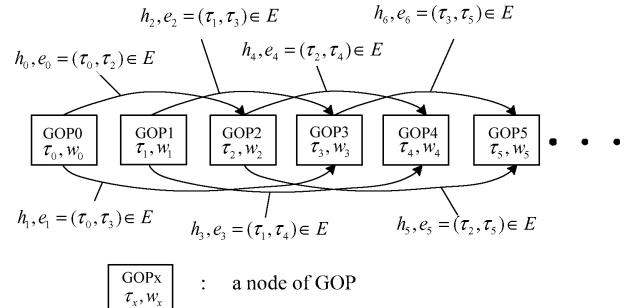


Fig. 4. DAG of MVC with two concurrent GOPs.

dependencies, and reflects the constraints and “cost” structure imposed by the target platform, and traversing through this DG reflects the performances of a particular scheduling algorithm.

C. Heuristics Algorithm Design Using DAG—Overview

We introduced two types of graphs in the previous section; one models the dynamic dependencies between encoding tasks among encoding units, and the other models the capabilities of the multicore platform.

To find a good heuristic algorithm q using these two graphs, we introduce another set of parameters X , which quantitatively describe the performance of the scheduling algorithm. Typical members of X may include the utilization of each core, the ratio of executing time to data transmission time, etc. Similarly, heuristics algorithm can also be parameterized by a vector Z , whose components may include parallel granularity, heuristics mechanism, etc. Z is a function of X , and each Z corresponds to a different q .

For example, when the utilization of core is low, a heuristics scheduling algorithm q may aim to generate a lower parallel granularity. To this end, q can lower the parallel granularity and choose to schedule a task that is prerequisite for the most of other tasks. If multiple tasks have the same dependent tasks, then the task with the shortest executing time should be scheduled.

On the other hand, when the ratio of executing time to data transmission is low, the goal of q maybe to reduce data transmission. To achieve this goal, q can schedule the tasks that are self-contained to each core thereby reducing data transmission.

D. Framework Design Using DAG

In this section, we demonstrate the design of heuristics mechanisms q using the DAG-based framework presented earlier. The framework (shown as Fig. 5) consists of three modules: the encoding module, the heuristics scheduling module and the multicore module.

Encoder Module: Video encoding algorithm is abstracted as a hierarchical multisource multisink DAG extended over the time axis. The DAG G is updated continuously to reflect the input that is currently “visible” to the encoder (e.g., the current frame and the frames in the reference frame buffer).

The main task of this module is to update the DAG once a task has been scheduled. $G(i)$ denotes the updated graph G after scheduling the i th task.

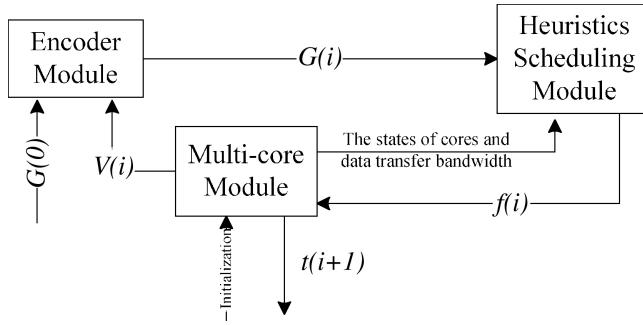


Fig. 5. Structure diagram of framework.

Heuristic Scheduling Module: The heuristics scheduling module synthesizes the states of encoder graph and multicore processor, and makes task scheduling decisions denoted by $f(i)$. Each heuristics mechanism corresponds to a particular scheduling algorithm.

Multicore Module: The module records and updates the states of the multicore processor, and calculates the set $V(i)$ of tasks which have been finished on the i th step.

Updates to the modules and graphs are only performed with i updated to $i + 1$ only when $V(i)$ is not empty.

Conceptually, the work flow is as the following.

- 1) **Initialization:** Initialize parameters and states of. Set $i = 0$, $G(i) = G(0)$, where $G(0)$ is the initial g (the number of concurrent GOPs) GOPs.
The states of cores and data transmission bandwidths are set free.
- 2) **Scheduling:** Based on the states of cores and data transmission and $G(i)$, q chooses the next scheduling action $f(i)$ based on both video encoding and processor constraints.
- 3) **Update multicore states:** When a task is scheduled, the states of the processor must be updated.
- 4) **Calculate the set V of nodes which are finished:** When a round of tasks for all the cores has been completely scheduled, the virtual time is incremented and the set $V(i)$ of tasks that have been completed during the previous time period is calculated. Otherwise, if the tasks in the current round have not been completely scheduled, the module will stall. If after action $f(i)$ is taken, no task can be scheduled in $G(i)$, or if there are no free cores, update virtual time t and $V(i)$, otherwise t cannot be updated and $V(i)$ is null. Here, the term “virtual time” refers to the total execution time as captured by the simulation framework (as opposed to measured on the actual platform).
- 5) **Update graph and simplify graph G :** When new nodes in current round have been encoded, i.e., $V(i)$ is not empty, graph G needs to be updated to $G(i + 1)$ referring $V(i)$. If $V(i)$ is empty, $G(i + 1) = G(i)$.
- 6) Repeat steps 2–5 until all frames are encoded.

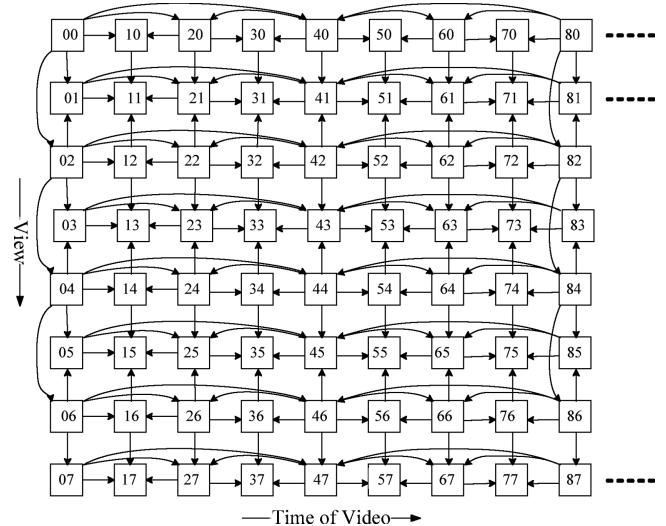


Fig. 6. GOP prediction structure in MVC.

IV. APPLICATION TO MVC ON SYMMETRIC MULTICORE ARCHITECTURES

To further illustrate the DAG-based approach, in this section, we use the scheduling of MVC on the Intel symmetric multicore processor as an example.

A. Framework for MVC

1) **Encoder Module for MVC:** In this particular example, we assume that the processing capability of each of the symmetric cores is c . This is a reasonable assumption because for each frame, the execution time of ME, DCT, quantization, and entropy coding is much higher than that of data transfer on symmetrical multicore architecture with shared memory, such as the Intel multicore processor. Under these two assumptions, graph $G(\Gamma, W, E, H)$ is simplified to $G(\Gamma, W/c, E)$ with no weighting for each edge. We also focus on the MVC encoder with a standard 8×8 GOP with the topological structure shown as in Fig. 6 [13]. Here, W/c is the execution time of tasks. We obtained W/c through experiment and categorize the tasks according to their, respectively, times. Through experiments we found that the ratio of the execution times for different tasks remain relatively constant for different clips, even though the value of the execution time itself may change.

During initialization, $G(0)$ is initialized to g connected GOPs. As defined earlier, Γ is the set of tasks in g GOPs, E is the set of edges of the encoder DAG. We use kd to denote the number of GOPs in the input.

Under these assumptions, pseudo code of the update procedure **UpdateG()** for G is shown in **UpdateG()**. In our experiments we chose frame as the parallel granularity, while GOP and MB are also possible choices. Through experiments we categorize the frames into five classes, namely I, P, B2, B3, and B4. I and P types correspond I-frame and P-frame, respectively, Bn ($n = 2, 3, 4$) type corresponds B-frame which refers to n frames. The estimated execution time of the node is estimated-based on the frame type. The most effective way of setting the weights is to base the value on the ratio of the time the processor needs for performing different tasks.

UpdateG() algorithm pseudo code:

```

 $\Gamma = \Gamma - V(i)$ 
//  $G(i)$  delete nodes of  $V(i)$  and relative edges
if ( $kd < g$ )
     $\Gamma = \Gamma + (g - kd)GOP$  //  $G(i)$  add the new nodes
    //and corresponding edges to make  $kd$  equal to  $g$ 
end if

```

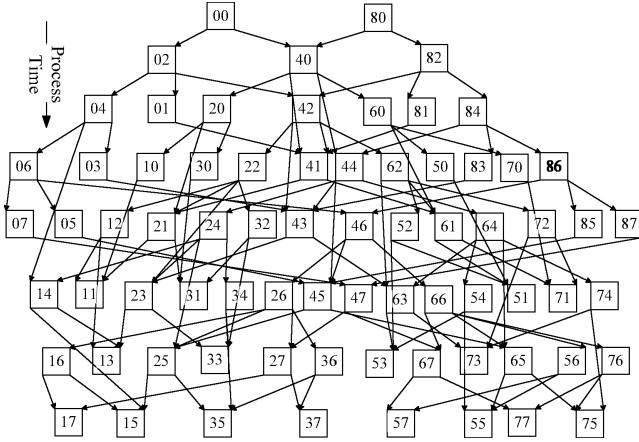


Fig. 7. GOP scheduling in MVC.

The ratio usually is relatively easier to obtain. In the paper, for example, to establish the framework for MVC, we need know the ratio of running times for encoding I-frame, P-frame and B – frames ($n = 2, 3, 4$). Such data can be estimated using virtually any implementation of MVC on the target platform. Fig. 7 shows the scheduling of a GOP in MVC. The scheduling (Fig. 7) can be generated from Fig. 6 automatically by topological sorting. The dependencies in Fig. 6 remain in Fig. 7. In order not to make Fig. 7 too complex, we eliminated some nonessential edges, e.g., edge from frame 00 to 20, the edges from 00 to 40 and from 40 to 20 imply that frame 20 depends on 00.

2) *Heuristics Scheduling Module for MVC*: In our experiments, we tested several different heuristic scheduling algorithms: Poc, Refs. Num., Ex. Time, Sum. Time, and Left Frame, which respectively gives the highest priority to: 1) the earliest frame in display order; 2) the frame which is referred by the most frames; 3) the frame with the shortest expected execution time; 4) on the subsequent route in G , the frame whose sum of executing time is the largest; and 5) on the subsequent route in G , the frame which is referred by the most frames. Task pool $TP(i)$ is the set of nodes which are ready to process in graph G . The task which has the highest priority is scheduled first. The scheduling action of the i th task to the $u(i)$ th core is denoted by $f(i) = (fn(i), u(i))$, $i = 0, 1, \dots, n - 1$, in which $fn(i)$ is the serial number of task.

3) *Multicore Module for MVC*: Because we assume that the cost of communication is negligible, the multicore module simply records and updates $S_{\text{core}} = (\text{core}_0, \text{core}_1, \dots, \text{core}_{m-1})$. The states of the cores is reflected in terms of remaining time for the task. We assume $\text{core}_j = (ct_j, cn_j)$, $j = 0, 1, \dots, m - 1$, ct_j denotes the left time of task scheduled on the core. If $ct_j = 0$, the corresponding core is free. cn_j denotes the serial number of task scheduled on the core, τ_{cn_j} is the

Update_MC() algorithm pseudo-code:

```

 $ct_{u(i)} = w_{fn(i)}/c$ 
 $cn_{u(i)} = fn(i)$ 
if (TP = null or  $\prod_{j=0}^{m-1} ct_j \neq 0$ )
    deltaT = min( $ct_j$ ,  $j = 0, 1, \dots, m - 1$ )
    t = t + deltaT //the virtual time goes forward
    for (j = 0 : m-1)
         $ct_j = ct_j - \Delta t$ 
        //the left working time reduces deltaT
        if ( $ct_j = 0$ )
             $V = V + \tau_{cn_j}$  //add encoded node to V
    end for
else
    V = null;
end if

```

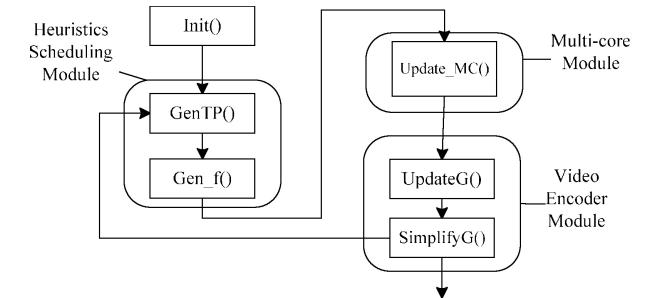


Fig. 8. Working flowchart of framework.

scheduled task. ct_j , cn_j , t and V are initialized as zero or null. The multicore update algorithm is shown in ***Update_MC()***. ***Update_MC()*** finds the minimal expected execution time Δt of the running cores, and emulates the execution by decreasing each anticipated running time (ct_i , $i = 0, 1, \dots, m - 1$) by Δt .

The flowchart of framework is in Fig. 8. ***Init()*** is initialization, ***GenTP()*** generates TP , ***Gen_f()*** generates action f , and ***SimplifyG()*** simplifies graph G .

B. Experimental Results

We implemented the simulation framework using C++ and a parallel MVC algorithm on the Intel multicore processor. We set quantitative parameter to 32, the number of views to 8, with each view containing 97 frames and the prediction structure as in the JMVC [13]. The NoC GOPs, heuristics scheduling algorithm, the number of cores and test sequences were configurable.

Figs. 9 and 11 show the speedup of five types of heuristics scheduling algorithms (Poc, Refs. Num., Ex. Time, Sum. Time and Left Frame) running on from one to 16 cores with two or three concurrent GOPs on the framework. Different criteria for scheduling the tasks were tested and shown in the figures.

As can be seen in Fig. 9, the slope of the speedup is linear when the NoC GOPs is 2 and the number of cores equal to or smaller than 6; when the number of cores is between 7 and 12, the slope of speedup lowers slightly.

Fig. 11 shows that Poc is the best heuristics scheduling algorithm, and slope of the speedup is linear when the NoC GOPs is 3 and the number of cores is increased to ten. The speedup slope lowers afterward.

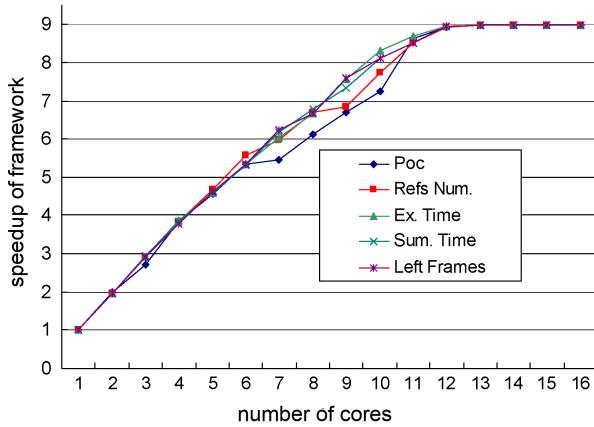


Fig. 9. Speedup ratio of five heuristics algorithms estimated by the proposed framework (NoC GOPs = 2).

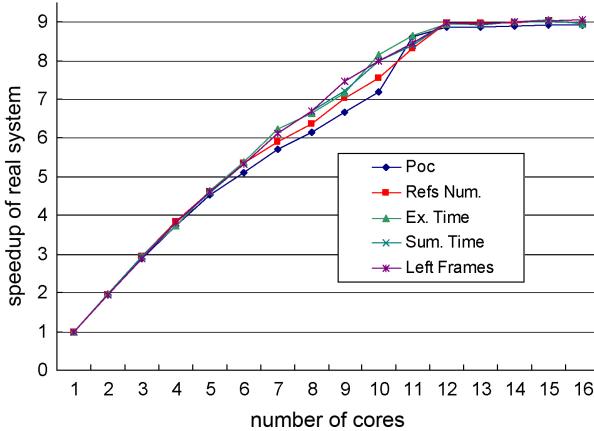


Fig. 10. Speedup ratio of five heuristics algorithms measured on the real system for test sequence *golf1* (NoC GOPs = 2).

We then tested the standard *golf1* sequence (320×240) [14] on the real system with the same parameters as Figs. 10 and 12. As can be seen from the figures, the actual performance on the actual system is nearly identical to the prediction in Figs. 9 and 11. This comparison shows that our framework can be used as a tool for reliably evaluating the performance of different scheduling algorithms without implementing the application and the scheduling algorithm on the actual platform.

Figs. 13 and 14 show the speedup of Poc encoder with 2, 3, 4 and an unlimited number of concurrent GOPs, as analyzed using the proposed framework and measured on the actual system, respectively. Again, the prediction and the actual measurement matches nicely. We can also see that as the NoC GOP increases, the speedup also increases. When the number of concurrent GOPs is unlimited, the slope of the speedup of dynamic heuristics algorithms is nearly linear.

We also tested other standard test sequences including *Rena* (640×480) [15], *Ballet* (1024×768) and *Breakdancer* (1024×768). The results are shown in Fig. 15.

It is clear from Fig. 15 that proposed framework and algorithms perform for these additional test sequences are

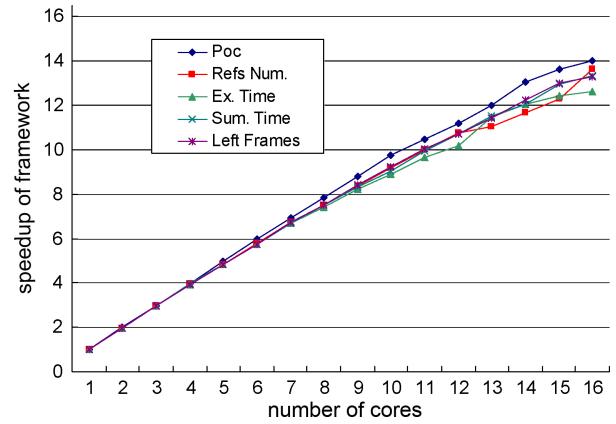


Fig. 11. Speedup ratio of five heuristics algorithms estimated by the proposed framework (NoC GOPs = 3).

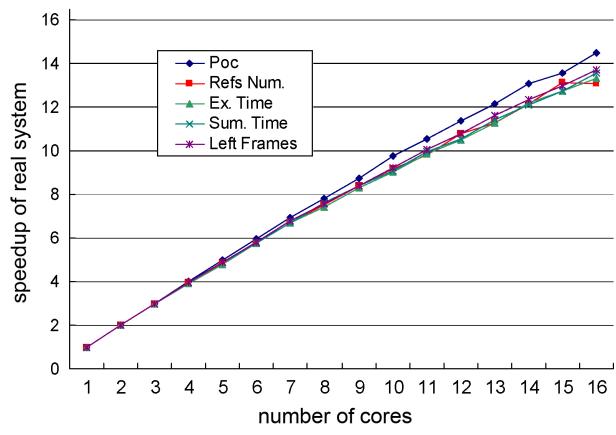


Fig. 12. Speedup ratio of five heuristics algorithms measured on the real system for test sequence *golf1* (NoC GOPs = 3).

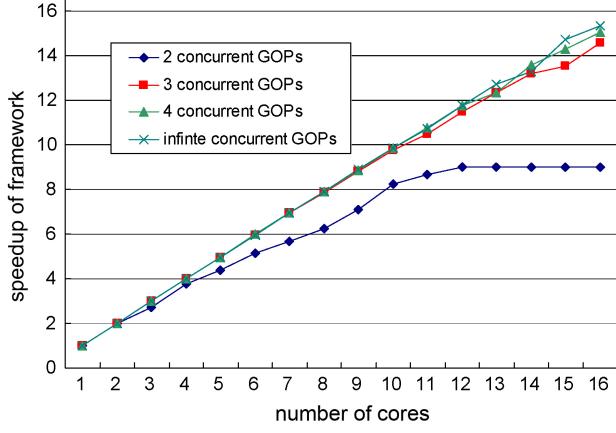


Fig. 13. Speedup ratio as a function of the number of concurrent GOPs, estimated by the proposed framework.

similar to that of *golf1*, the speedup is linear with respect to the number of cores and the analysis with the framework matches the actual measurement nicely.

In summary, the comparison of the framework and performance on the real system shows that our framework can be used to predict and analyze critical characteristics of the performance of an algorithm for an application on the actual

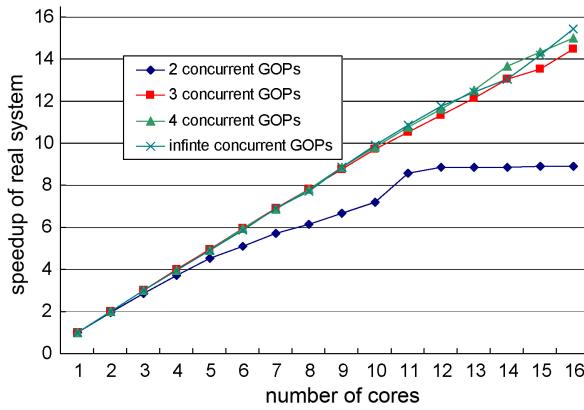


Fig. 14. Speedup Ratio as a function of the number of concurrent GOPs, measured on the real system for test sequence *golfI*.

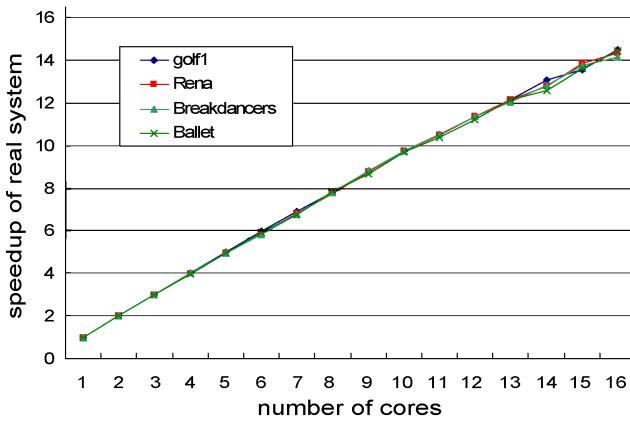


Fig. 15. Speedup ratio for different test sequences measured on the real system (NoC GOPs=3).

system. As the analysis using the framework is computationally trivial, while implementing applications and scheduling algorithms on the actual platform is not only time consuming and error prone, the framework presented in this paper can serve as an important tool in the design and analysis for building applications on multicore platforms. Because of the reasonable parameters abstraction of the critical characteristics in MVC, the framework matches the real system very well, as can be seen from the experimental results. It also validates that our previous assumption that the cost of communication is negligible is reasonable.

In this paper, to better explain the core idea of the proposed framework, we use frame as the granularity for parallelism and five widely-used heuristic-based scheduling algorithms. The significant difference in performance of the five algorithms (all of which make intuitive sense) show that even for this simple example, the problem of finding a good dynamic scheduling algorithm is still highly challenging. This further proves the necessity for using a low complexity and generic framework for predicting the performance of the scheduling algorithm.

V. CONCLUSION AND FUTURE RESEARCH

The increased interests in the design and adaption of a multicore platform for computationally intensive and dynamic applications such as MVCs necessitate dynamic scheduling algorithms and tools for analysis and fine tuning such algorithms, and for predicting important characteristics of the performance (e.g., core usage, speedup ratio, bandwidth, etc.).

In this paper, a framework for modeling the behavior of the target application and target platform using DAGs was introduced, with simulation results demonstrating the usefulness and scalability of the framework and methodology presented.

The framework and tools introduced in this paper are generic, expandable to other applications and optimality metrics, and lightweight. It can be used to accurately capture the performance of applications and algorithms that would otherwise take much longer to implement, debug, and test, and/or for platforms that are not yet available for actual implementation.

In our experiments for this paper, the scheduling algorithms that we tested for each test scenario are static, i.e., even though input to the algorithm maybe highly dynamic, the criteria used by the algorithm for tasking scheduling (e.g., Poc or otherwise) remains unchanged for the entire application.

Because of the dynamic and lightweight real-time nature of the tools that we introduced, a perceivable and important area for future research is embedding the tools introduced in the paper into the application with feedbacks from the physical platform (e.g., core utilization, etc.) incorporated into the decision loop so that a truly dynamic and real-time optimized scheduling can be achieved.

ACKNOWLEDGMENT

The authors would like to thank Intel China Research Center for providing an experimental environment.

REFERENCES

- [1] K. Asanovic, R. Bodik, B. C. Catanzaro, J. J. Gebis, P. Husbands, K. Keutzer, D. A. Patterson, W. L. Plishker, J. Shalf, S. W. Williams, and K. A. Yelick, "The landscape of parallel computing research: A view from Berkeley," Dept. Electrical Eng. Comput. Sci., Univ. California, Berkeley, Tech. Rep. UCB/EECS-2006-183, 2006.
- [2] K. Shen and E. J. Delp "A parallel implementation of an MPEG1 encoder: Faster than real-time!" in *Proc. Soc. Photo-Optical Instrum. Eng. Conf. Digital Video Compression: Algorithms Tech.*, vol. 2419. Bellingham, WA, 1995, pp. 407–418.
- [3] H. C. Yung and K. K. Leung, "Spatial and Temporal Data Parallelization of the H.261 video coding algorithm," *IEEE Trans. Circuits Syst. Video Technol.*, vol. 11, no. 1, pp. 91–104, Jan. 2001.
- [4] B. T. Erik and G. J. Egbert "Mapping of MPEG-4 decoding on a flexible architecture platform," in *Proc. Soc. Photo-Optical Instrum. Eng. Media Processor*, vol. 4674. 2002, pp. 1–13.
- [5] Y. K. Chen, E. Q. Li, X. Zhou, and S. Ge, "Implementation of H.264 encoder and decoder on personal computers," *J. Visual Commun. Image Representation*, vol. 17, no. 2, pp. 509–532, Apr. 2006
- [6] H. Wang, X. Mao, and L. Yu "A novel HDTV decoder and decentralized control scheme," *IEEE Trans. Consum. Electron.*, vol. 47, no. 4, pp. 723–728, Nov. 2001.
- [7] A. K. Nanda, J. R. Moulic, R. E. Hanson, G. Goldrian, M. N. Day, B. D. Amora, and S. Kesavarapu, "Cell/B.E. blades: Building blocks for scalable, real-time, interactive, and digital media servers," *IBM J. Res. Dev.*, vol. 51, no. 5, pp. 573–582, Sep. 2007.

- [8] Y. Yang, G. Jiang, M. Yu, and D. Zhu "Parallel process of hyper-space-based multiview video compression," in *Proc. IEEE Int. Conf. Image Process.*, Atlanta, GA, Oct. 2006, pp. 521–524.
- [9] Y. Pang, L. F. Sun, S. L. Guo, and S. Q. Yang, "Spatial and temporal data parallelization of multiview video encoding algorithm," in *Proc. IEEE Int. Workshop Multimedia Signal Process.*, Crete, Greece, Oct. 2007, pp. 441–444.
- [10] I. Ahmad, Y. Kwok, and M. Wu, "Performance comparison of algorithms for static scheduling of DAGs to multiprocessors," in *Proc. 2nd Aust. Conf. Parallel Real-Time Syst.*, Sep. 1995, pp. 185–192.
- [11] A. A. Steen, M. David, D. P. Michael, and M. T. Christopher, "On the relation between conditional independence models determined by finite distributive lattices and by directed acyclic graphs," *J. Stat. Plan. Inference*, vol. 48, no. 1, pp. 25–46, Nov. 1995.
- [12] Q. S. Hua and Z. G. Chen, "Efficient granularity and clustering of the directed acyclic graphs," in *Proc. 4th Int. Conf. Parallel Distributed Comput., Applicat. Technol.*, Aug. 2003, pp. 625–628.
- [13] *Joint Multiview Video Model (JMVM) 8.0*, Document MPEG N9762.doc, MPEG, Archamps, France, Apr. 2008.
- [14] *KDDI Multiview Video Sequences for MPEG 3DAV Use*, Document M10533.doc, MPEG, Munich, Mar. 2004.
- [15] *Test Sequences with Different Camera Arrangements for Call for Proposals on Multiview Video Coding*, Document M12338.doc, MPEG, Poznan, Poland, Jul. 2005



Yi Pang (S'07) received the B.S. and M.S. degrees in computer science from Tsinghua University, Beijing, China, in 2005 and 2007, respectively. She is currently pursuing the Ph.D. degree from the Department of Computer Science and Technology, Tsinghua University.

During her B.S. and Ph.D. study, she held several internships: Summer Senior Analyst at Lehman Brothers, Hong Kong, a Software Engineer at Betaresearch, Munich, Germany, an Instructor at Qinghai University, Xining, China, etc. Her research interests include parallelism of multimedia (e.g., multiview video coding, H.264, ray tracing, etc.) on multicore processors (e.g., Cell/B.E., Intel multicore, etc.).



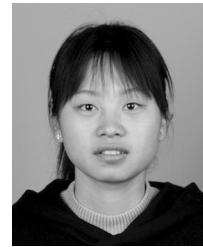
Lifeng Sun (M'05) received the B.S., M.S., and Ph.D. degrees in system engineering from the National University of Defense Technology, Hunan, China, in 1995, 1997, and 2000, respectively.

He is currently an Associated Professor in the Department of Computer Science and Technology, Tsinghua University, Beijing, China. His research interests include interactive multiview video, video sensor networks, peer-to-peer streaming, and distributed video coding.

Jiangtao (Gene) Wen (SM'05) received the B.S., M.S., and Ph.D. degrees (with honors), all in electrical engineering, from Tsinghua University, Beijing, China, in 1992, 1994, and 1996, respectively.

From 1996 to 1998, he was a Staff Research Fellow at the University of California, Los Angeles (UCLA), where he conducted cutting-edge research on multimedia coding and communications. Many of his inventions there were later adopted by international standards such as H.263, MPEG, and H.264. After UCLA, he served as the Principal Scientist at PacketVideo Corp., the Chief Technical Officer at Morphibus Technology Inc., the Director of Video Codec Technologies at Mobilygen Corp., and as a Technology Advisor at Ortiva Wireless and Stretch, Inc. Since 2009, he has been a Professor at the Department of Computer Science and Technology, Tsinghua University. He is a world-renowned expert in multimedia communication over hostile networks, video coding, and communications. He has authored many widely referenced papers in related fields. Products deploying technologies that he developed are currently widely used worldwide. He holds over 30 patents with numerous others pending.

Dr. Wen is an Associate Editor for the IEEE TRANSACTIONS ON CIRCUITS AND SYSTEMS FOR VIDEO TECHNOLOGY.



Fengyan Zhang was born in 1983. She received the B.S. degree in computer science and technology from Tsinghua University, Beijing, China, in 2006. She is currently pursuing the M.S. degree in computer science and technology from the Department of Computer Science and Technology, Tsinghua University.

Her current research interests include parallel video coding and multiview video coding, among others.



Weidong Hu was born in 1986. He received the B.S. degree in computer science and technology from Tsinghua University, Beijing, China, in 2009. He is currently pursuing the M.S. and Ph.D. degrees in computer science and technology from the Department of Computer Science and Technology, Tsinghua University.

His current research interests include parallelism of multimedia on multicore processors.

Mr. Hu won the gold medal from the Association for Computing Machinery International Collegiate Programming Contest in March 2007, and the gold medal in the International Olympiad in Informatics (IOI) in 2004 and in 2005. He is the coach of the IOI China Team.



Wei Feng was born in 1987. He is currently pursuing the B.S. degree in computer science and technology from Tsinghua University, Beijing, China.

His research interests include high-performance computing, parallel video coding, and multiview video coding, among others.



Shiqiang Yang (M'97–SM'08) graduated from the Department of Computer Science and Technology, Tsinghua University, Beijing, China, in 1977, and received the M.E. degree in 1983.

He is now a Professor at Tsinghua University. His research interests include multimedia technology and systems, video compression and streaming, content-based retrieval, semantics for multimedia information, and embedded multimedia systems. He has published more than 100 papers in international conferences and journals.

Mr. Yang is currently the President of the Multimedia Committee of the China Computer Federations.